

1. Problem statement

The assignment is to build an **automatic bin-level SKU verification system** for a warehouse / e-commerce setting.

For each storage bin we have:

- A **bin image** (`bin-images/XXXX.jpg`)
- A **metadata JSON** (`metadata/XXXX.json`) listing:
 - ASIN / SKU
 - Product title
 - Expected quantity in that bin

The system must:

1. Detect all products visible in the bin image.
2. Match detections to the SKUs from metadata.
3. Estimate per-SKU quantity from detections.
4. Compare predicted vs expected quantities and classify each SKU as:
 - `FULL_MATCH` – correct or nearly correct count
 - `PARTIAL_OCCLUDED` – under-count, likely due to occlusion
 - `NOT_FOUND` – SKU expected but not detected
 - `OVERCOUNT` – more items detected than expected (potential false positives)

On top of this, the project includes:

- Full **EDA** on images, metadata and model outputs.
- A sequence of **pipeline versions** (v1–v5) with a comparison study.
- A small **Streamlit app** that behaves like a “smart invoice / order validator” on top of the model outputs.

2. Data and repository structure

Relevant top-level layout (after your refactor):

```
applied-ai-assignment1/
├── app/
│   ├── __init__.py
│   └── streamlit_app.py          # Streamlit UI (main app
entrypoint)
```

```
|─ bin-images/                                # Bin photos (00001.jpg,  
00002.jpg, ...)  
|  └ ... (data, usually not all in GitHub)  
  
|─ legacy_pipelines/  
|  ├ __init__.py  
|  ├ validate_full_pipeline_v2.py  # Older baseline pipeline  
|  └ validate_full_pipeline_v4_clip018.py  
  
|─ metadata/                                 # Per-bin JSON metadata  
(00001.json, ...)  
  
|─ mlops/  
|  ├ __init__.py  
|  ├ evaluate_model_v3_counts.py    # Evaluation + count-style  
metrics  
|  └ plot_results.py                      # CLI plotting for metrics  
  
|─ models/  
|  ├ __init__.py  
|  └ validate_full_pipeline_v5.py  # Final CV pipeline (v5:  
BOX=0.17, CLIP=0.20)  
  
|─ notebooks/                               # EDA / experimentation notebooks  
  
|─ out/  
|  ├ summary_*.csv                         # Evaluation summaries (v1–v5)  
|  ├ *_detected_*.jpg                     # Visualisations with drawn boxes  
|  └ inference_log.csv (optional)  
  
|─ weights/  
|  ├ GroundingDINO_SwinT_OGC.py          # GroundingDINO config  
|  └ groundingdino_swinT_ogc.pth          # GroundingDINO weights (large,  
may be local only)  
  
|─ .gitignore  
|─ Assignment #1 [Applied AI for Industry] (1).pdf
```

```
└─ README.md
└─ evaluate_model.py          # Older eval script
└─ evaluate_model_v2.py
└─ plot_results.py           # Older plotting script
└─ requirements.txt
└─ s3_downloader.py          # Helper to fetch data from S3
└─ sanity_check.py           # Quick sanity run script
└─ validate_full_pipeline.py # Very first strict prototype
└─ validate_full_pipeline_v3_dino_loose.py # Intermediate "loose
DINO" version
```

- Each **JSON** has a `BIN_FCSKU_DATA` dict mapping ASIN → `{name, quantity}`.

The **summary CSVs** are one row per (bin, SKU) with:

```
image_id, asin, sku_name, expected, visible, status[, pretty_status]
```

- - `expected` – quantity from metadata
 - `visible` – quantity inferred from detections
 - `status` – categorical evaluation label

Dataset size (for the main experiments):

- **~987–988 bins**
 - **2702–2704 SKU-bin pairs** depending on version (v2–v5).
-

3. Exploratory Data Analysis (EDA)

You implemented an `EDA.py` script that:

- Flattens all metadata JSONs into a single table
- Analyses image characteristics
- Analyses all `summary_*.csv` files
- Writes plots and CSVs into `out/eda/`

3.1 Metadata EDA

From the flattened metadata:

- Each row: `image_id`, `asin`, `sku_name`, `expected`
- Key findings (qualitative, from the script's outputs):
 - **SKUs per bin**: a moderate number (typical small bin with a handful of SKUs).
 - **Expected quantities**: vary from 1 to higher counts; most rows have small integers (1–6).
 - Some SKUs appear in **multiple bins** (popular products) – captured by `value_counts` on `asin`.

Plots produced:

- Histogram of **SKUs per bin** – helps understand workload per image.
- Histogram of **expected quantity per row** – shows whether the model must count many instances or just presence.

These support design choices:

- Because quantities are small, a **hard cap** on detections per SKU (expected + tolerance) is reasonable.
- Many bins have multiple SKUs, so a **multi-prompt** text-conditioned detector (GroundingDINO) fits well.

3.2 Image EDA

The script scans `bin-images` with Pillow and logs:

- Width, height and aspect ratio (w / h)
- Orientation (landscape / portrait / square)

Typical findings:

- Most images are similar resolution; aspect ratios lie around a narrow band.
- A mix of landscape and portrait bins, which justifies using **normalized bounding boxes** from GroundingDINO.

Plots:

- Histograms of **width**, **height**, **aspect ratio**
- Bar chart of **orientation counts**

These ensure there's no systematic anomaly (e.g. a subset of very tiny or corrupted images).

3.3 Summary / Model Output EDA

For each `summary_v*.csv` the EDA script prints:

- Total SKU rows and #bins
- Distribution of `status`
- Presence metrics treating “visible > 0” as “SKU detected”
- Scatter plots of `expected` vs `visible` and status bar charts
- An aggregated table `summary_presence_metrics.csv` comparing versions.

This EDA is the backbone for the comparison study in Section 6.

4. Methodology

4.1 Model choices

The system combines:

1. **GroundingDINO** (Swin-T OGC variant)
 - Text-conditioned object detector.
 - Takes natural language prompts of product names.
 - Produces bounding boxes with text scores per region.
2. **OpenAI CLIP** (`ViT-B/32` variant)
 - Image–text embedding model.
 - Used as a second-stage verifier:
 - Crop each candidate box.
 - Compute similarity between crop and SKU text.
 - Reject boxes with low CLIP score.

This two-stage design is important:

- GroundingDINO is strong at *localising* objects based on language but may fire on wrong regions when prompts are long or cluttered.
- CLIP is strong at judging **global image–text alignment** and filters many false positives.

4.2 Per-bin pipeline

All `validate_full_pipeline_v*.py` versions follow the same high-level structure:

1. **Load metadata** for a bin (`image_id.json`).
 - Extract list of SKUs: (`asin`, `name`, `expected_qty`).
2. **Prepare text prompts** for GroundingDINO.
 - Typically the product title, possibly shortened / aliased.
3. **Run GroundingDINO** with:

- `BOX_THRESHOLD` – minimum detector box score.
 - `TEXT_THRESHOLD` – minimum text relevance.
 - `IOU_THRESHOLD_PER_SKU` – NMS IOU for per-SKU de-duplication (0.55 in all versions).
4. **Run CLIP verification:**
- For each candidate box and SKU text, compute cosine similarity.
 - Keep those with similarity $\geq \text{CLIP_THRESHOLD}$.
5. **Count per SKU:**
- Group remaining boxes by ASIN.
 - Apply any version-specific caps / tolerance.
 - Compute `visible` count for each SKU.

Classify status for each SKU row:

```
if visible == 0:
    status = "NOT_FOUND"
elif abs(visible - expected) <= 1:
    status = "FULL_MATCH"
elif visible < expected:
    status = "PARTIAL_OCCLUDED"
else:
    status = "OVERCOUNT"
```

6.

7. **Write summary row and visualisation image:**

- One CSV row per SKU with `status` and `pretty_status`.
- `.jpg` with bounding boxes and labels.

5. Pipeline versions

5.1 Version overview


The main versions and their hyper-parameters:



Version file	Label (infor mal)	BOX_THRE SHOLD	TEXT_THRE SHOLD	CLIP_THRE SHOLD	Notes
--------------	-------------------------	-------------------	--------------------	--------------------	-------

<code>validate_full_pipeline.py</code>	v1 – strict	0.18	0.18	0.22	Very conservative, only run on 1 bin for sanity check
<code>validate_full_pipeline_v2.py</code>	v2 – baseline	0.18	0.18	0.20	Slightly relaxed CLIP; full dataset
<code>validate_full_pipeline_v3_dino_loose.py</code>	v3 – loose DINO	0.15	0.15	0.20	More permissive detector → more boxes
<code>validate_full_pipeline_v4_clip018.py</code>	v4 – loose CLIP	0.15	0.15	0.18	Very permissive CLIP → many more candidates
<code>validate_full_pipeline_v5.py</code>	v5 – final count-aware	0.17	0.18	0.20	Midpoint DINO, stricter CLIP, hard cap per SKU

Key design evolution:

- **v1 → v2**: relax CLIP from 0.22 to 0.20, keeping DINO thresholds; this was the first usable baseline.
- **v2 → v3**: loosen GroundingDINO (0.18 → 0.15) to increase recall; CLIP fixed at 0.20.
- **v3 → v4**: further loosen CLIP to 0.18 to maximise recall, at the cost of more false positives.
- **v4 → v5**:
 - Tighten DINO slightly to 0.17.
 - Restore CLIP to 0.20.
 - Introduce a **count-aware cap**: for each SKU, keep at most `expected + COUNT_TOLERANCE` boxes (with `COUNT_TOLERANCE = 1`).
 - Use the `status_from_counts` helper with ± 1 tolerance baked into `FULL_MATCH`.





6. Quantitative comparison of versions

Below we compare v2–v5 on the full dataset. v1 is omitted from tables because it only contains **3 rows / 1 bin** (debug run) and is not representative.

6.1 Status distributions

For each summary we count how many SKU rows fall into each status:

Summary file	FULL_MATCH	PARTIAL_OCCLUDE D	OVERCOUNT	NOT_FOUND
v2_relaxed	1311 (48.5%)	256 (9.5%)	196 (7.2%)	941 (34.8%)
v3_dino_loose	1396 (51.6%)	262 (9.7%)	319 (11.8%)	727 (26.9%)
v4_clip018	1617 (59.8%)	271 (10.0%)	629 (23.3%)	185 (6.8%)
v5_clip020_box0 17	1699 (62.9%)	174 (6.4%)	0 (0.0%)	829 (30.7%)

Observations:

- Moving from **v2 → v3 → v4**:
 - **FULL_MATCH** steadily increases (48.5% → 51.6% → 59.8%).
 - **NOT_FOUND** dramatically drops (34.8% → 26.9% → 6.8%) – recall is much higher.
 - But **OVERCOUNT** explodes (7.2% → 11.8% → 23.3%) as we relax thresholds.
- **v5**:
 - Achieves the **highest FULL_MATCH rate (62.9%)**.
 - Completely removes **OVERCOUNT** thanks to the count cap (**expected + 1**).
 - However **NOT_FOUND** climbs back to ~30.7% because we tightened detection / CLIP.

6.2 Detection recall and count error

Define:

- **Detection recall**: proportion of SKU rows with **visible > 0** (i.e. model detected at least one instance).
- **Mean absolute error (MAE)**: average of **abs(visible - expected)** across all rows.

Summary:

Summary file	Rows	Bins	Detection recall	FULL_MATCH rate	NOT_FOUND rate	PARTIAL rate	OVERCOUNT rate	MAE (visible - expected)
v2_relaxed	2704	988	0.652	0.485	0.348	0.095	0.072	1.49
v3_dino_loose	2704	988	0.731	0.516	0.269	0.097	0.118	1.51
v4_clip018	2702	987	0.932	0.598	0.068	0.100	0.233	1.59
v5_clip020_box017	2702	987	0.693	0.629	0.307	0.064	0.000	1.37

Key trade-offs:

1. **v2 → v3** (looser DINO)
 - Detection recall increases (65% → 73%).
 - FULL_MATCH improves modestly (48.5% → 51.6%).
 - NOT_FOUND falls by ~8 percentage points.
 - OVERCOUNT rises from 7% to 12%.
 - MAE slightly worsens (1.49 → 1.51).
2. Interpretation: looser DINO helps find more SKUs but also introduces more false boxes, especially when prompts are generic or products are visually similar.
3. **v3 → v4** (looser CLIP)
 - Detection recall jumps to **93.2%** – almost every SKU has at least one detection.
 - NOT_FOUND becomes very rare (~7%).
 - FULL_MATCH improves further to ~60%.
 - However OVERCOUNT balloons to **23.3%**, nearly a quarter of all rows.
 - MAE is worst (1.59).
4. Interpretation: CLIP at 0.18 allows many borderline matches; the model “sees” almost everything but over-counts aggressively, often counting background clutter as additional units.
5. **v4 → v5** (count-aware, mid thresholds)
 - DINO threshold raised to 0.17; CLIP back to 0.20; and a **count cap** of **expected + 1**.

- Detection recall drops relative to v4 (93% → 69%), roughly back to between v2 and v3.
 - FULL_MATCH climbs to **62.9%**, the best among all versions.
 - PARTIAL_OCCLUDED shrinks to **6.4%**, lowest across versions.
 - OVERCOUNT becomes **0% by design**: the cap prevents heavy over-counting.
 - MAE improves from 1.59 (v4) to **1.37**, the best among v2–v5.
6. Interpretation: v5 intentionally sacrifices some recall to get much more reliable counts, avoiding scary overestimates. For an inventory/audit scenario, under-counting due to occlusion can be handled by human review, while over-counting can directly lead to wrong shipping or stock records.

6.3 Role of the count-aware cap

In v5, for each SKU the algorithm:

1. Sorts candidate boxes by CLIP similarity.
2. Limits accepted boxes to `min(num_candidates, expected + 1)`.

As a result:

- **Overcount by more than 1** is impossible.
- In practice, the dataset has many cases where model tends to over-fire slightly. The cap ensures those extra boxes are dropped.
- Because `FULL_MATCH` is defined with ± 1 tolerance, a mild over-count (`expected=3, visible=4`) is still considered `FULL_MATCH`, which is acceptable for inventory scenarios.

This change is the main reason why **OVERCOUNT goes from 23.3% (v4) to 0% (v5)** even though CLIP threshold is only slightly stricter.

7. Streamlit app: “Smart Bin Order Validator”

The `streamlit_app.py` provides an interactive demo of the pipeline:

7.1 Pipeline selection

The app defines a `PIPELINES` dictionary:

```
PIPELINES = {
    "v2 - BOX=0.18, CLIP=0.20 (baseline)": {...},
```

```
"v4 - BOX=0.18, CLIP=0.18 (relaxed CLIP)": {...},  
"v5 - BOX=0.17, CLIP=0.20 (count-aware)": {...},  
}
```

In the **sidebar**, the user can:

- Choose the pipeline version (`v2`, `v4`, or `v5`).
- Choose which **bin ID** (image) to run on.
- Click a button to run the selected model.

The app uses cached functions:

- `@st.cache_resource` to load models (GroundingDINO + CLIP) per version.
- `@st.cache_data` to cache metadata / summary results.

7.2 Bin-level view

Once a bin is processed:

- The app shows the original **bin image** (with or without drawn boxes, depending on implementation).
- It presents a table (DataFrame) with columns like:
 - `asin`, `sku_name`, `expected`, `visible`, `status`

This is basically a visual front-end over the `summary_v*.csv` outputs for a single bin.

7.3 Order / invoice validation

The more interesting part is the “invoice” layer:

1. Function `build_order_from_bin(df_bin: pd.DataFrame)`:
 - Extracts unique SKUs for that bin with `expected` and `visible`.
 - For each row, the UI shows:
 - Product name + ASIN
 - A `number_input("Order_qty")` where the user specifies how many units they want to `order`.

Builds an order table:

```
asin, sku_name, order_qty, expected, visible
```

◦

2. Line-item validation:

For each line, a simple rule:

```
if visible >= order_qty:  
    order_status = "PASS"  
else:  
    order_status = "FAIL"  
  
○  
○ The app shows an order table with order_status and an order-level summary:  
    “Order-level summary: X/Y line items are PASS.”
```

This turns the bin verification pipeline into a **smart order validator**:

- You visually see what's in the bin.
- You “place” an order based on that bin.
- The model checks whether the bin actually contains enough stock to fulfil the order.

This nicely links the ML pipeline to a real business use-case (inventory picking / order verification).

8. Implementation details

8.1 Requirements

`requirements.txt` contains:

```
torch  
torchvision  
Pillow  
numpy  
pandas  
opencv-python  
tqdm  
boto3  
transformers  
groundingdino  
git+https://github.com/openai/CLIP.git
```

Additional packages used in EDA/app (not all in the base file yet):

- `matplotlib` – EDA plots
- `streamlit` – app
- (optional) `seaborn`, `jupyter` – for notebooks

8.2 Scripts

- `s3_downloader.py` – retrieve images/metadata from S3.
- `validate_full_pipeline_v*.py` – core per-bin pipelines.
- `evaluate_model*.py` / `evaluate_model_v3_counts.py` – compute metrics and comparisons from summary files.
- `plot_results*.py` – CLI scripts to plot status distributions and trends.
- `EDA.py` – dataset + output analysis, as described above.

The project is fully **script-driven**, which makes experiments reproducible: you can regenerate each summary CSV by re-running the corresponding `validate_full_pipeline_*` script.

9. Discussion and insights

9.1 What worked well

- **GroundingDINO + CLIP** is a powerful combination for SKU-level detection in noisy bins.
- Careful tuning of **BOX and CLIP thresholds** showed a clear trade-off:
 - Lower thresholds → higher recall but more overcounts.
 - Higher thresholds → fewer false positives but more missed items.
- Introducing a **count-aware cap** in v5 dramatically improved practical usability:
 - Eliminated overcounting.
 - Achieved the highest rate of `FULL_MATCH` (62.9%).
 - Lowered mean absolute count error to 1.37.

9.2 Remaining issues

- **Occlusion** is still a big challenge:
 - Many `PARTIAL_OCCLUDED` cases come from items behind others, partially cut off, or obstructed by packaging.
- **Fine-grained product differences**:
 - Some SKUs differ only by minor text/colour on packaging; GroundingDINO + CLIP may struggle to distinguish them.
- The current evaluation does **not** explicitly track false positives on *completely unexpected SKUs*, because the summaries are keyed by metadata SKUs only. Extra SKUs not in metadata are currently visible only in visualisations and not quantified.

9.3 Possible future improvements

- Use **instance segmentation** or **depth cues** to better handle overlapping items.
 - Train a small **SKU-level classifier** on top of cropped boxes to disambiguate similar packages.
 - Introduce a **calibration step** where thresholds are optimised on a held-out validation set using metrics like F1 for presence and MAE for counts.
 - Extend evaluation to include **unexpected SKUs** by logging detections that did not match any metadata entry.
-

10. Conclusion

This assignment delivers an end-to-end system that:

1. Ingests bin images and structured metadata.
2. Uses **GroundingDINO** and **CLIP** to detect and verify products.
3. Produces per-SKU counts and status labels summarised in CSVs.
4. Iteratively improves model performance across **five pipeline versions**, culminating in **v5**, which balances recall, accurate counting, and safety against overcounting.
5. Provides both:
 - An **EDA toolkit** for understanding dataset and outputs.
 - A **Streamlit “Smart Bin Order Validator” app** that turns model outputs into a practical decision tool for checking whether a bin can fulfil an order.

In terms of comparison, **v4** achieves the highest raw recall but at the cost of many overcounts and larger count errors, while **v5** offers the best overall trade-off for real-world inventory.